

Dataset Information

The objective of this task is to detect hate speech in tweets. For the sake of simplicity, we say a tweet contains hate speech if it has a racist or sexist sentiment associated with it. So, the task is to classify racist or sexist tweets from other tweets.

Formally, given a training sample of tweets and labels, where label '1' denotes the tweet is racist/sexist and label '0' denotes the tweet is not racist/sexist, your objective is to predict the labels on the test dataset.

For training the models, we provide a labelled dataset of 31,962 tweets. The dataset is provided in the form of a csv file with each line storing a tweet id, its label and the tweet.



Import modules

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import re
import string
import nltk
import warnings
%matplotlib inline

warnings.filterwarnings('ignore')
```

Loading the dataset

```
In [2]: df = pd.read_csv('Twitter Sentiments.csv')
df.head()
```

```
Out[2]:
```

	id	label	tweet
0	1	0	@user when a father is dysfunctional and is s...
1	2	0	@user @user thanks for #lyft credit i can't us...
2	3	0	bihday your majesty
3	4	0	#model i love u take with u all the time in ...
4	5	0	factsguide: society now #motivation

```
In [3]: # datatype info
df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31962 entries, 0 to 31961
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0    id      31962 non-null    int64
1    label    31962 non-null    int64
2    tweet    31962 non-null    object
dtypes: int64(2), object(1)
memory usage: 749.2+ KB
```

Preprocessing the dataset

```
In [4]: # removes pattern in the input text
def remove_pattern(input_txt, pattern):
    r = re.findall(pattern, input_txt)
    for word in r:
        input_txt = re.sub(word, "", input_txt)
    return input_txt
```

```
In [5]: df.head()
```

```
Out[5]:
```

	id	label	tweet
0	1	0	@user when a father is dysfunctional and is s...
1	2	0	@user @user thanks for #lyft credit i can't us...
2	3	0	bihday your majesty
3	4	0	#model i love u take with u all the time in ...
4	5	0	factsguide: society now #motivation

```
In [6]: # remove twitter handles (@user)
df['clean_tweet'] = np.vectorize(remove_pattern)(df['tweet'], "@[\w]*")
```

```
In [7]: df.head()
```

```
Out[7]:
```

	id	label	tweet	clean_tweet
0	1	0	@user when a father is dysfunctional and is s...	when a father is dysfunctional and is so sel...
1	2	0	@user @user thanks for #lyft credit i can't us...	thanks for #lyft credit i can't use cause th...
2	3	0	bihday your majesty	bihday your majesty
3	4	0	#model i love u take with u all the time in ...	#model i love u take with u all the time in ...
4	5	0	factsguide: society now #motivation	factsguide: society now #motivation

```
In [8]: # remove special characters, numbers and punctuations
df['clean_tweet'] = df['clean_tweet'].str.replace("[^a-zA-Z#]", " ")
df.head()
```

```
Out[8]:
```

	id	label	tweet	clean_tweet
0	1	0	@user when a father is dysfunctional and is s...	when a father is dysfunctional and is so sel...
1	2	0	@user @user thanks for #lyft credit i can't us...	thanks for #lyft credit i can t use cause th...
2	3	0	bihday your majesty	bihday your majesty
3	4	0	#model i love u take with u all the time in ...	#model i love u take with u all the time in ...
4	5	0	factsguide: society now #motivation	factsguide society now #motivation

```
In [9]: # remove short words
df['clean_tweet'] = df['clean_tweet'].apply(lambda x: " ".join([w for w in x.split() if len(w) > 3]))
df.head()
```

```
Out[9]:
```

	id	label	tweet	clean_tweet
0	1	0	@user when a father is dysfunctional and is s...	when father dysfunctional selfish drags kids i...
1	2	0	@user @user thanks for #lyft credit i can't us...	thanks #lyft credit cause they offer wheelchai...
2	3	0	bihday your majesty	bihday your majesty
3	4	0	#model i love u take with u all the time in ...	#model love take with time
4	5	0	factsguide: society now #motivation	factsguide society #motivation

```
In [10]: # individual words considered as tokens
tokenized_tweet = df['clean_tweet'].apply(lambda x: x.split())
tokenized_tweet.head()
```

```
Out[10]:
```

0	[when, father, dysfunctional, selfish, drags, ...]
1	[thanks, #lyft, credit, cause, they, offer, wh...]
2	[bihday, your, majesty]
3	[#model, love, take, with, time]
4	[factsguide, society, #motivation]

Name: clean_tweet, dtype: object

```
In [11]: # stem the words
from nltk.stem.porter import PorterStemmer
stemmer = PorterStemmer()

tokenized_tweet = tokenized_tweet.apply(lambda sentence: [stemmer.stem(word) for word in sentence])
tokenized_tweet.head()
```

```
Out[11]:
```

0	[when, father, dysfunct, selfish, drag, kid, i...]
1	[thank, #lyft, credit, caus, they, offer, whee...]
2	[bihday, your, majesti]
3	[#model, love, take, with, time]
4	[factsguid, societi, #motiv]

Name: clean_tweet, dtype: object

```
In [12]: # combine words into single sentence
for i in range(len(tokenized_tweet)):
    tokenized_tweet[i] = " ".join(tokenized_tweet[i])

df['clean_tweet'] = tokenized_tweet
df.head()
```

Out[12]:	id	label	tweet	clean_tweet
0	1	0	@user when a father is dysfunctional and is s...	when father dysfunct selfish drag kid into dys...
1	2	0	@user @user thanks for #lyft credit i can't us...	thank #lyft credit caus they offer wheelchair ...
2	3	0		bihday your majesty
3	4	0	#model i love u take with u all the time in ...	#model love take with time
4	5	0	factsguide: society now #motivation	factsguid societi #motiv

Exploratory Data Analysis

In [14]: `# !pip install wordcloud`

Collecting wordcloud

Downloading wordcloud-1.8.1-cp38-cp38-win_amd64.whl (155 kB)

Requirement already satisfied: pillow in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (7.2.0)

Requirement already satisfied: numpy>=1.6.1 in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (1.18.5)

Requirement already satisfied: matplotlib in c:\programdata\anaconda3\lib\site-packages (from wordcloud) (3.2.2)

Requirement already satisfied: python-dateutil>=2.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (2.8.1)

Requirement already satisfied: cyclor>=0.10 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (0.10.0)

Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (2.4.7)

Requirement already satisfied: kiwisolver>=1.0.1 in c:\programdata\anaconda3\lib\site-packages (from matplotlib->wordcloud) (1.2.0)

Requirement already satisfied: six>=1.5 in c:\programdata\anaconda3\lib\site-packages (from python-dateutil>=2.1->matplotlib->wordcloud) (1.15.0)

Installing collected packages: wordcloud

Successfully installed wordcloud-1.8.1

In [15]: `# visualize the frequent words`

`all_words = " ".join([sentence for sentence in df['clean_tweet']])`

`from wordcloud import WordCloud`

`wordcloud = WordCloud(width=800, height=500, random_state=42, max_font_size=100).generate`

`# plot the graph`

`plt.figure(figsize=(15,8))`

`plt.imshow(wordcloud, interpolation='bilinear')`

`plt.axis('off')`


`plt.show()`

Out[24]:

	Hashtag	Count
0	run	72
1	lyft	2
2	disapoint	1
3	getthank	2
4	model	375

In [25]:

```
# select top 10 hashtags
d = d.nlargest(columns='Count', n=10)
plt.figure(figsize=(15,9))
sns.barplot(data=d, x='Hashtag', y='Count')
plt.show()
```



In [26]:


```
freq = nltk.FreqDist(ht_negative)
d = pd.DataFrame({'Hashtag': list(freq.keys()),
                  'Count': list(freq.values())})
d.head()
```

Out[26]:

	Hashtag	Count
0	cnn	10
1	michigan	2
2	tcot	14
3	australia	6
4	opkillingbay	5

In [27]:

```
# select top 10 hashtags
d = d.nlargest(columns='Count', n=10)
plt.figure(figsize=(15,9))
sns.barplot(data=d, x='Hashtag', y='Count')
```

Input Split

```
In [28]: # feature extraction
from sklearn.feature_extraction.text import CountVectorizer
bow_vectorizer = CountVectorizer(max_df=0.90, min_df=2, max_features=1000, stop_words='e
bow = bow_vectorizer.fit_transform(df['clean_tweet'])
```

```
In [35]: # bow[0].toarray()
```

```
In [36]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(bow, df['label'], random_state=42, t
```

Model Training

```
In [41]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import f1_score, accuracy_score
```

```
In [38]: # training
model = LogisticRegression()
model.fit(x_train, y_train)
```

```
Out[38]: LogisticRegression()
```

```
In [39]: # testing
pred = model.predict(x_test)
f1_score(y_test, pred)
```

```
Out[39]: 0.49763033175355453
```

```
In [42]: accuracy_score(y_test, pred)
```

Out[42]: 0.9469403078463271

```
In [43]: # use probability to get output  
pred_prob = model.predict_proba(x_test)  
pred = pred_prob[:, 1] >= 0.3  
pred = pred.astype(np.int)  
  
f1_score(y_test, pred)
```

Out[43]: 0.5545722713864307

```
In [44]: accuracy_score(y_test, pred)
```

Out[44]: 0.9433112251282693

```
In [47]: pred_prob[0][1] >= 0.3
```

Out[47]: False

```
In [ ]:
```